

# TritNet

Optimising matrix multiply computation  
through scaling bit-wise logical operations

Ryan Cherian

August 2024

## Contents

<b>1</b>	<b>Axioms of logical operations</b>	<b>2</b>
1.1	Scalar logical algebra . . . . .	2
1.2	Multi-linear logical algebra . . . . .	2
<b>2</b>	<b>Lemmas to represent multiplication of non-bit sets as logical operations</b>	<b>3</b>
<b>3</b>	<b>Axioms of matrix multiplication</b>	<b>4</b>
<b>4</b>	<b>Matrix multiplication using only bitwise operations</b>	<b>5</b>
4.1	Multiplication of binary matrices in $\{0, 1\}$ . . . . .	5
4.2	Multiplication of binary matrices in $\{-1, 1\}$ . . . . .	5
4.3	Multiplication of trinary matrices in $\{-1, 0, 1\}$ . . . . .	5
4.4	Multiplication of trinary-binary matrices . . . . .	6
<b>5</b>	<b>Extension</b>	<b>7</b>
<b>6</b>	<b>Consequences</b>	<b>8</b>
6.1	Problem formulation . . . . .	8
6.2	Options . . . . .	8

# 1 Axioms of logical operations

## 1.1 Scalar logical algebra

Note that throughout this derivation, we abide the notation that Greek letters represent bits, and all other variables (incl. binary variables outside  $\{0, 1\}$ , e.g. in  $\{-1, 1\}$ ) are Roman. The axioms in sum are:

$$\forall \alpha, \beta \in \{0, 1\}$$

**Axiom 1.1** (AND):

$$\alpha \wedge \beta = \alpha\beta$$

**Axiom 1.2** (OR):

$$\alpha \vee \beta = \alpha + \beta - \alpha\beta$$

**Axiom 1.3** (NOT):

$$\bar{\alpha} = 1 - \alpha$$

$\alpha$	$\beta$	$\alpha\beta$
0	0	0
0	1	0
1	0	0
1	1	1

Table 1: AND

Table 1: Truth tables

$\alpha$	$\beta$	$\alpha + \beta - \alpha\beta$
0	0	0
0	1	1
1	0	1
1	1	1

Table 2: OR

$\alpha$	$1 - \alpha$
0	1
1	0

Table 1: NOT

## 1.2 Multi-linear logical algebra

A population count (`_popc()`) is equivalent to a *full* tensor contraction (one form of multi-linear reduction) in integer or floating-point arithmetic - i.e. the sum of all elements of an n-rank tensor.

**Axiom 1.4** (Population Count):

$$\text{popc}(\alpha) := \sum_i \alpha_i$$

Thus, if a tensor can be represented in first-rank (vector) form, and as a concatenation of strictly Boolean operations, it's full contraction can be represented as a pop. count. For instance:

$$\begin{aligned} \sum_{i,j} \alpha_{i,j} \vee \beta_{i,j} &= \text{popc}(\alpha \vee \beta) \\ \sum_k \alpha_{i,k} \wedge \beta_{k,j} &= \text{popc}(\alpha_i \wedge \beta_j) = \alpha_i \cdot \beta_j \end{aligned}$$

Where the last equality represents summations of the dot product of Boolean vectors. In other words, the *matrix multiplication* of Boolean matrices.

## 2 Lemmas to represent multiplication of non-bit sets as logical operations

**Lemma 1** (OR):  $\alpha \vee \beta = \alpha\beta! + \beta$

*Proof.*

$$\begin{aligned}\alpha \vee \beta &= \alpha + \beta - \alpha\beta \\ &= \alpha(1 - \beta) + \beta = \alpha\beta! + \beta\end{aligned}$$

□

**Lemma 2** (XOR):  $\alpha \veebar \beta = \alpha\beta! + \beta\alpha! = \alpha + \beta - 2\alpha\beta$

*Proof.*

$$\begin{aligned}\alpha \veebar \beta &= \alpha\beta! + \beta\alpha! \quad (\text{truth table}) \\ &= \alpha(1 - \beta) + \beta(1 - \alpha) = \alpha + \beta - 2\alpha\beta\end{aligned}$$

□

**Lemma 3** (Transform):  $a \in \{-1, 1\} \mapsto \alpha \in \{1, 0\}$  through  $\alpha = \frac{1-a}{2}$

*Proof.* Simple linear transformation:

$$\forall a \in \{-1, 1\} \exists \alpha = \frac{1-a}{2} : \alpha \in \{1, 0\}$$

□

**Lemmas 4.x** Multiplication of binary variables:

**Lemma 4.0** Binary multiplication in  $\alpha, \beta \in \{0, 1\}$  :  $\alpha\beta = \alpha \wedge \beta$

(This is Axiom 1)

**Lemma 4.1** Binary multiplication of  $\forall a, b \in \{-1, 1\}$  :  $ab = 1 - 2(\alpha \veebar \beta)$

*Proof.*

$$\begin{aligned}\textbf{Lemma 3} \implies ab &= (1 - 2\alpha)(1 - 2\beta) = 1 - 2\alpha - 2\beta + 4\alpha\beta \\ &= 1 - 2(\alpha + \beta - 2\alpha\beta) \\ \textbf{Lemma 2} \implies &= 1 - 2(\alpha \veebar \beta)\end{aligned}$$

□

**Lemma 4.2:** Binary multiplication  $\alpha \in \{0, 1\}$  with  $b \in \{-1, 1\}$  :  $\alpha b = \alpha - 2\alpha\beta = \alpha \veebar \beta - \beta$

*Proof.*

$$\begin{aligned}\textbf{Lemma 3} \implies ab &= \alpha(1 - 2\beta) \quad \forall \alpha \in \{0, 1\}, b \in \{-1, 1\} \\ &= \alpha - 2\alpha\beta = \alpha - 2(\alpha \wedge \beta) \\ \textbf{Lemma 2} \implies &= \alpha \veebar \beta - \beta\end{aligned}$$

□

### 3 Axioms of matrix multiplication

General matrix multiplication  $\mathbf{Z} = \mathbf{WA}$  is defined as follows:

$$\mathbf{Axiom\ 2.1:} \quad \exists \mathbf{Z} = \mathbf{WA} : Z_{i,j} = \sum_k W_{i,k} A_{k,j} \quad \forall W_{i,k}, A_{k,j}, Z_{i,j} \in \mathbb{R}$$

If using NVIDIA Tensor Cores, we can define two new binary matrix multiplication types:

$$\mathbf{Axiom\ 2.2:} \quad \exists \mathbf{Z} = \mathbf{W} * \mathbf{\alpha} : Z_{i,j} = \sum_k (\omega_{i,k} \wedge \alpha_{k,j}) \quad \forall \omega_{i,k}, \alpha_{k,j} \in \{0, 1\}, Z_{i,j} \in [0..k]$$

$$\mathbf{Axiom\ 2.3:} \quad \exists \mathbf{Z} = \mathbf{W} \star \mathbf{\alpha} : Z_{i,j} = \sum_k (\omega_{i,k} \vee \alpha_{k,j}) \quad \forall \omega_{i,k}, \alpha_{k,j} \in \{0, 1\}, Z_{i,j} \in [0..k]$$

Note that NVIDIA's cores each operate on a maximum of (8, 128, 8) dimensions.

## 4 Matrix multiplication using only bitwise operations

### 4.1 Multiplication of binary matrices in $\{0, 1\}$

$$\begin{aligned}
 Z_{i,j} &= \sum_k \omega_{i,k} \alpha_{k,j} & \forall \omega_{i,k}, \alpha_{k,j} \in \{0, 1\} \\
 \text{Axiom 1.1} \implies Z_{i,j} &= \sum_k \omega_{i,k} \wedge \alpha_{k,j} \\
 \text{Axioms 2.1, 2.2} \implies \therefore \mathbf{Z} &= \mathbf{\omega} \mathbf{\alpha} = \mathbf{\omega} * \mathbf{\alpha}
 \end{aligned}$$

### 4.2 Multiplication of binary matrices in $\{-1, 1\}$

$$\begin{aligned}
 Z_{i,j} &= \sum_k W_{i,k} A_{k,j} & \forall W_{i,k}, A_{k,j} \in \{-1, 1\}, Z_{i,j} \in k \cdot [W \times A] \\
 \text{Lemma 4.1} \implies &= \sum_k 1 - 2(\omega_{i,k} \vee \alpha_{k,j}) \\
 &= k - 2 \sum_k (\omega_{i,k} \vee \alpha_{k,j}) \\
 \text{Axioms 2.1, 2.3} \implies \therefore \mathbf{Z} &= \mathbf{W} \mathbf{A} = \mathbf{K} - 2(\mathbf{\omega} \star \mathbf{\alpha}) & \forall K_{i,j} \in \{k\}
 \end{aligned}$$

So for a  $(m, k, n)$  matrix multiply, calculated compactly through bitwise operations on  $N$ -bit words:

```

Z[i][j] += N - 2*_popc(W[i][k] ^ A[k][j]);

//or equivalently
Z[i][j] += N - (_popc(W[i][k] ^ A[k][j])<<1);

//or when utilising NVIDIA Tensor Core capability:
Z[i][j] = n*k - 2*(bmma_PopXor(W,A))[i][j]; // _popc(W[i][k] ^ A[k][j])

```

### 4.3 Multiplication of trinary matrices in $\{-1, 0, 1\}$

$$\begin{aligned}
 W_{i,j}, A_{i,j} &\in \{-1, 0, 1\} \forall i, j \in \mathbb{N} : \mathbf{Z} = \mathbf{W} \mathbf{A} \\
 \exists \omega_{i,j}^o, \alpha_{i,j}^o &\in \{0, 1\}, W'_{i,j}, A'_{i,j} \in \{-1, 1\} : \mathbf{W} = \mathbf{\omega}^o \odot \mathbf{W}', \mathbf{A} = \mathbf{\alpha}^o \odot \mathbf{A}' \\
 \therefore \mathbf{Z} &= (\mathbf{\omega}^o \odot \mathbf{W}')(\mathbf{\alpha}^o \odot \mathbf{A}') \\
 \text{Axiom 1} \implies Z_{i,j} &= \sum_k W_{i,k} A_{k,j} = \sum_k \omega_{i,k}^o W'_{i,k} \alpha_{k,j}^o A'_{k,j} \\
 &= \sum_k \omega_{i,k}^o \alpha_{k,j}^o W'_{i,k} A'_{k,j} \\
 \text{Lemma 4} \implies &= \sum_k \omega_{i,k}^o \alpha_{k,j}^o (1 - 2(\omega'_{i,k} \vee \alpha'_{k,j})) \\
 &= \sum_k \omega_{i,k}^o \alpha_{k,j}^o - 2 \sum_k \omega_{i,k}^o \alpha_{k,j}^o (\omega'_{i,k} \vee \alpha'_{k,j})
 \end{aligned}$$

This gives the calculation for each element. Thus you can stop here for the operation per kernel thread:

```
Z[i][j] += _popc(W0[i][k] & A0[k][j]) - 2*_popc(W0[i][k] & A0[k][j] & (W1[i][k] ^ A1[k][j]));
```

Or if utilising Tensor cores, continue by expanding the XOR operation using Lemma 2, then simplify:

$$\begin{aligned}
\text{Lemma 2} \implies &= \sum_k \omega_{i,k}^o \alpha_{k,j}^o - 2 \sum_k \omega_{i,k}^o \alpha_{k,j}^o (\overline{\omega'_{i,k}} \alpha'_{k,j} + \omega'_{i,k} \overline{\alpha'_{k,j}}) \\
&= \sum_k \omega_{i,k}^o \alpha_{k,j}^o - 2 \sum_k \omega_{i,k}^o \alpha_{k,j}^o \overline{\omega'_{i,k}} \alpha'_{k,j} - 2 \sum_k \omega_{i,k}^o \alpha_{k,j}^o \omega'_{i,k} \overline{\alpha'_{k,j}} \\
&= \sum_k \omega_{i,k}^o \alpha_{k,j}^o - 2 \sum_k (\omega_{i,k}^o \overline{\omega'_{i,k}}) (\alpha_{k,j}^o \alpha'_{k,j}) - 2 \sum_k (\omega_{i,k}^o \omega'_{i,k}) (\alpha_{k,j}^o \overline{\alpha'_{k,j}})
\end{aligned}$$

$$\begin{aligned}
\text{Axioms 1,2} \implies \therefore \mathbf{Z} = \mathbf{WA} &= \omega^o \alpha^o - 2(\omega^o \odot \overline{\omega'}) (\alpha^o \odot \alpha') - 2(\omega^o \odot \omega') (\alpha^o \odot \overline{\alpha'}) \\
&= \omega^o \alpha^o - 2\mathbf{PQ} - 2\mathbf{RS}
\end{aligned}$$

This translates to:

```
Z = bmma_PopAND(W0,A0) - 2*(bmma_PopAND(W0(!W1), A0A1)) - 2*(bmma_PopAND(W0W1, A0(!A1)));
```

#### 4.4 Multiplication of trinary-binary matrices

$$\begin{aligned}
W_{i,j} &\in \{-1, 0, 1\}, A_{i,j} \in \{-1, 1\} \forall i, j \in \mathbb{N} : \mathbf{Z} = \mathbf{WA} \\
\exists \omega_{i,j}^+, \omega_{i,j}^- &\in \{0, 1\} : \mathbf{W} = \omega^+ - \omega^-
\end{aligned}$$

$$\begin{aligned}
&\therefore \mathbf{Z} = (\omega^+ - \omega^-) \mathbf{A} \\
\text{Axiom 1} \implies &Z_{i,j} = \sum_k W_{i,k} A_{k,j} = \sum_k (\omega_{i,k}^+ - \omega_{i,k}^-) A_{k,j} \\
&= \sum_k \omega_{i,k}^+ A_{k,j} - \sum_k \omega_{i,k}^- A_{k,j} \\
\text{Lemma 4} \implies &= \sum_k (\omega_{i,k}^+ \vee \alpha_{k,j} - \alpha_{k,j}) - \sum_k (\omega_{i,k}^- \vee \alpha_{k,j} - \alpha_{k,j}) \\
&= \sum_k \omega_{i,k}^+ \vee \alpha_{k,j} - \sum_k \omega_{i,k}^- \vee \alpha_{k,j} \\
\text{Axiom 2.3} \implies &\therefore \mathbf{Z} = \mathbf{WA} = (\omega^+ \star \alpha) - (\omega^- \star \alpha)
\end{aligned}$$

Thus you can calculate each element using the following operation per kernel thread:

```
Z[i][j] += _popc(W+[i][k] ^ A[k][j]) - _popc(W-[i][k] ^ A[k][j]);
```

Or calculate the whole matrix using tensor cores, with the following:

```
Z = bmma_PopXOR(W+,A) - bmma_PopXOR(W-,A);
```

## 5 Extension

Note that this method is **extendable**. Consider the following split:

$$\begin{aligned}\forall \alpha &\in \{2, 1, 0, -1\} \\ \exists abcd = \alpha &\quad \forall a \in \{0, 1\}, b \in \{1, 1\}, c \in \{1, 2\}, d \in \{1, -1\}\end{aligned}$$

With this we can then simplify the multiplication of the two quaternary variables like so:

$$abcd * pqrs = ap * bq * cr * ds = \dots$$

However, we then require 4 bits to hold each quaternary variable - which could easily have been held in 2 bits. So this method is not feasible beyond multiplications of sets with 3 elements.